

Volatility models for REMX: GARCH, GARCH-X, GARCHAND, GARCHND and EGARCH-X

Summary of all models estimated in the notebooks **03-GP-models-gaussian**, **03-GP-models-student** and **04-extension-models**. Each section contains the variance expression, a brief description of what the model is meant to demonstrate, and the main estimation results (Gaussian first, Student-t second). The *Sig.* column summarises significance: *** = 1%, ** = 5%, * = 10%. A final part covers the six robustness checks from notebooks **06-robustness-checks** and **07-extra-robustness-checks**.

1. GARCH(1,1) — Baseline

Conditional variance:

$$\sigma_t^2 = \omega + \alpha \cdot r_{t-1}^2 + \beta \cdot \sigma_{t-1}^2$$

What we want to demonstrate: Canonical Bollerslev (1986) specification. It serves as the benchmark against which every sentiment-augmented model is compared. The goal is to confirm persistence ($\alpha + \beta$ close to 1) and that the basic parameters are significant; all later extensions are evaluated relative to this fit.

Gaussian

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0857	0.0301	2.8497	0.0044	***
alpha[1]	0.0893	0.0161	5.5536	0.0000	***
beta[1]	0.8975	0.0182	49.2043	0.0000	***

Log-likelihood : -5982.3069

AIC : 11970.6138

BIC : 11988.3882

alpha[1] + beta[1] = 0.9867 (stationary (<1))

Student-t

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0878	0.0299	2.9308	0.0034	***
alpha[1]	0.0841	0.0150	5.5880	0.0000	***
beta[1]	0.9010	0.0178	50.6845	0.0000	***
nu	11.1674	2.3239	4.8054	0.0000	***

Log-likelihood : -5956.2330

AIC : 11920.4660

BIC : 11944.1652

alpha[1] + beta[1] = 0.9850 (stationary (<1))

nu = 11.1674 (smaller -> heavier tails; Gaussian as nu -> infinity)

2. GARCH-X (tone) — G&P Eq. (4)

Conditional variance:

$$\sigma_t^2 = \omega + \alpha \cdot r_{t-1}^2 + \beta \cdot \sigma_{t-1}^2 + \gamma \cdot \text{tone}_{t-1}^2$$

What we want to demonstrate: Symmetric extension that injects the daily average news tone as an exogenous regressor in the variance equation. We want to test whether news intensity (regardless of sign) generates additional volatility in REMX, i.e. whether γ is statistically significant.

Gaussian

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0835	0.0457	1.8258	0.0679	*
alpha	0.0905	0.0159	5.7033	0.0000	***
beta	0.8961	0.0179	50.0241	0.0000	***
gamma	0.0024	0.0213	0.1136	0.9096	

Log-likelihood : -5982.5250

AIC : 11973.0499

BIC : 11996.7491

alpha + beta : 0.9867 (stationary (<1))

Student-t

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.1202	0.0460	2.6149	0.0089	***
alpha	0.0843	0.0153	5.5026	0.0000	***
beta	0.8984	0.0188	47.8456	0.0000	***
gamma	-0.0179	0.0155	-1.1549	0.2481	
nu	10.8940	2.2264	4.8932	0.0000	***

Log-likelihood : -5955.3892

AIC : 11920.7783

BIC : 11950.4023

alpha + beta : 0.9828 (stationary (<1))

Estimated degrees of freedom: nu_hat = 10.894

Note: t-statistic on nu vs 0 is not economically meaningful (nu > 2 by construction).

3. GARCH-X (log article growth) — G&P Eq. (4)

Conditional variance:

$$\sigma_t^2 = \omega + \alpha \cdot r_{t-1}^2 + \beta \cdot \sigma_{t-1}^2 + \gamma \cdot \log_art_growth_{t-1}^2$$

What we want to demonstrate: Same structure as the previous GARCH-X, but using the log-growth in article volume ($\log(1+N_t) - \log(1+N_{t-1})$) as a proxy for news intensity. The log transform replaces the raw growth rate to remove a single-day outlier (art_growth = 61 on 2015-10-23) that would otherwise dominate the squared term. We want to see whether sustained changes in media coverage (a significant γ) anticipate next-day volatility.

Gaussian

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0757	0.0362	2.0947	0.0362	**
alpha	0.0910	0.0160	5.6763	0.0000	***
beta	0.8959	0.0181	49.5268	0.0000	***
gamma	0.0203	0.0372	0.5476	0.5840	

Log-likelihood : -5982.2558

AIC : 11972.5115

BIC : 11996.2107

alpha + beta : 0.9869 (stationary (<1))

Student-t

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0885	0.0359	2.4669	0.0136	**
alpha	0.0852	0.0150	5.6883	0.0000	***
beta	0.8996	0.0176	51.0511	0.0000	***
gamma	0.0012	0.0356	0.0346	0.9724	
nu	11.1206	2.3100	4.8142	0.0000	***

Log-likelihood : -5956.3176

AIC : 11922.6352

BIC : 11952.2592

alpha + beta : 0.9848 (stationary (<1))

Estimated degrees of freedom: nu_hat = 11.121

Note: t-statistic on nu vs 0 is not economically meaningful (nu > 2 by construction).

4. GARCHAND (tone) — G&P Eq. (5)

Conditional variance:

$$\sigma_t^2 = \omega + \alpha \cdot r_{t-1}^2 + \beta \cdot \sigma_{t-1}^2 + \gamma (1 + d_1 \cdot \theta) \cdot \text{tone}_{t-1}^2, d_1 = 1\{\text{tone}_{t-1} < 0\}$$

What we want to demonstrate: Asymmetric version of GARCH-X: when the previous day's tone is negative, its impact on variance is amplified by the factor $(1 + \theta)$. We want to demonstrate that negative news drives a stronger volatility response than positive news ($\theta > 0$ and significant).

Gaussian

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0899	0.0353	2.5503	0.0108	**
alpha	0.0904	0.0159	5.6758	0.0000	***
beta	0.8958	0.0180	49.7159	0.0000	***
gamma	-0.0009	0.0041	-0.2274	0.8201	
theta	0.4979	6.7407	0.0739	0.9411	

Log-likelihood : -5982.5410

AIC : 11975.0819

BIC : 12004.7059

alpha + beta : 0.9862 (stationary (<1))

Fraction of days with tone_{t-1} < 0 (d1 = 1): 0.481

Student-t

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.1284	0.0475	2.7003	0.0069	***
alpha	0.0839	0.0159	5.2868	0.0000	***
beta	0.8966	0.0201	44.5784	0.0000	***
gamma	-0.0005	0.0002	-2.5306	0.0114	**
theta	47.9057	16.0878	2.9778	0.0029	***
nu	10.9218	2.2666	4.8185	0.0000	***

Log-likelihood : -5954.5020

AIC : 11921.0040

BIC : 11956.5528

alpha + beta : 0.9805 (stationary (<1))

Fraction of days with tone_{t-1} < 0 (d1 = 1): 0.481

Estimated degrees of freedom: nu_hat = 10.922

5. GARCHAND (log article growth) — G&P Eq. (6)

Conditional variance:

$$\sigma_t^2 = \omega + \alpha \cdot r_{t-1}^2 + \beta \cdot \sigma_{t-1}^2 + \gamma \cdot d_2 \cdot \log_art_growth_{t-1}^2, d_2 = 1\{\log_art_growth_{t-1} > 0\}$$

What we want to demonstrate: The article-volume effect is switched on only on days in which the news flow GROWS relative to the previous day ($\log_art_growth > 0$). We want to show that novelty (rising coverage) is what moves volatility, rather than low news volumes.

Gaussian

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0730	0.0356	2.0533	0.0400	**
alpha	0.0903	0.0159	5.6646	0.0000	***
beta	0.8964	0.0181	49.4520	0.0000	***
gamma	0.0521	0.0736	0.7077	0.4791	

Log-likelihood : -5982.0589

AIC : 11972.1179

BIC : 11995.8171

alpha + beta : 0.9867 (stationary (<1))

Fraction of days with $\log_art_growth_{\{t-1\}} > 0$ ($d_2 = 1$): 0.476

Student-t

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0842	0.0351	2.4029	0.0163	**
alpha	0.0851	0.0149	5.7110	0.0000	***
beta	0.8998	0.0176	51.0085	0.0000	***
gamma	0.0177	0.0677	0.2610	0.7941	
nu	11.1560	2.3245	4.7994	0.0000	***

Log-likelihood : -5956.2750

AIC : 11922.5499

BIC : 11952.1739

alpha + beta : 0.9849 (stationary (<1))

Fraction of days with $\log_art_growth_{\{t-1\}} > 0$ ($d_2 = 1$): 0.476

Estimated degrees of freedom: $\nu_{\text{hat}} = 11.156$

6. GARCHND (tone & log article growth) — G&P Eq. (7)

Conditional variance:

$$\sigma_t^2 = \omega + \alpha \cdot r_{t-1}^2 + \beta \cdot \sigma_{t-1}^2 + \gamma \cdot d_3 \cdot x_{t-1}^2, d_3 = 1\{\sigma_{t-1}^2 \geq \kappa\}, x \in \{\text{tone}, \text{log_art_growth}\}$$

What we want to demonstrate: The news impact activates only when the previous day's variance exceeds a threshold κ , calibrated from an annualised volatility level (30% and 50%). We want to test whether the sentiment effect is non-linear in the volatility regime: a significant γ in the high-volatility regime would imply that news amplifies volatility only when the market is already stressed.

Gaussian

GARCHND x = Tone kappa = 30% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0967	0.0247	3.9157	0.0001	***
alpha	0.0888	0.0101	8.7465	0.0000	***
beta	0.8891	0.0139	63.7458	0.0000	***
gamma	0.0375	0.0342	1.0965	0.2729	

Log-likelihood : -5979.6159

AIC : 11967.2319

BIC : 11990.9311

alpha + beta : 0.9779 (stationary (<1))

d3 active in-sample ($\sigma_{t-1}^2 \geq \kappa$): 0.616

GARCHND x = Tone kappa = 50% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0532	0.0115	4.6104	0.0000	***
alpha	0.0872	0.0027	32.6900	0.0000	***
beta	0.9101	0.0032	287.2016	0.0000	***
gamma	-0.4811	0.0227	-21.2193	0.0000	***

Log-likelihood : -5977.0528

AIC : 11962.1055

BIC : 11985.8047

alpha + beta : 0.9973 (stationary (<1))

d3 active in-sample ($\sigma_{t-1}^2 \geq \kappa$): 0.072

GARCHND x = log_art_growth kappa = 30% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.1088	0.0216	5.0458	0.0000	***
alpha	0.0945	0.0309	3.0598	0.0022	***
beta	0.8795	0.0116	75.8273	0.0000	***
gamma	0.1299	0.4061	0.3200	0.7490	

Log-likelihood : -5979.8694

AIC : 11967.7387

BIC : 11991.4379

alpha + beta : 0.9740 (stationary (<1))

d3 active in-sample ($\sigma^2_{t-1} \geq \kappa$): 0.621

GARCHND x = log_art_growth kappa = 50% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0551	0.0134	4.1159	0.0000	***
alpha	0.0901	0.0014	65.0193	0.0000	***
beta	0.9077	0.0003	2997.4667	0.0000	***
gamma	-1.6273	0.0741	-21.9478	0.0000	***

Log-likelihood : -5975.4760

AIC : 11958.9520

BIC : 11982.6512

alpha + beta : 0.9978 (stationary (<1))

d3 active in-sample ($\sigma^2_{t-1} \geq \kappa$): 0.066

Student-t

GARCHND x = Tone kappa = 30% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0918	0.0282	3.2498	0.0012	***
alpha	0.0850	0.0137	6.2072	0.0000	***
beta	0.8981	0.0162	55.5481	0.0000	***
gamma	0.0081	0.0239	0.3400	0.7339	
nu	11.2155	2.2412	5.0041	0.0000	***

Log-likelihood : -5956.2451

AIC : 11922.4903

BIC : 11952.1143

alpha + beta : 0.9830 (stationary (<1))

d3 active in-sample ($\sigma^2_{t-1} \geq \kappa$): 0.613

Estimated degrees of freedom: $\nu_{\text{hat}} = 11.215$

GARCHND x = Tone kappa = 50% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0481	0.0174	2.7651	0.0057	***
alpha	0.0810	0.0047	17.3331	0.0000	***
beta	0.9165	0.0078	116.7886	0.0000	***
gamma	-0.7553	0.0470	-16.0863	0.0000	***
nu	11.5147	2.3839	4.8302	0.0000	***

Log-likelihood : -5952.0998

AIC : 11914.1995

BIC : 11943.8235

alpha + beta : 0.9975 (stationary (<1))
d3 active in-sample ($\sigma^2_{t-1} \geq \kappa$): 0.062
Estimated degrees of freedom: $\nu_{\text{hat}} = 11.515$

GARCHND x = log_art_growth kappa = 30% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.1010	0.0299	3.3781	0.0007	***
alpha	0.0877	0.0126	6.9575	0.0000	***
beta	0.8915	0.0163	54.5380	0.0000	***
gamma	0.0620	0.0480	1.2913	0.1966	
nu	11.2835	2.2238	5.0740	0.0000	***

Log-likelihood : -5955.7018
AIC : 11921.4036
BIC : 11951.0276
alpha + beta : 0.9793 (stationary (<1))
d3 active in-sample ($\sigma^2_{t-1} \geq \kappa$): 0.621
Estimated degrees of freedom: $\nu_{\text{hat}} = 11.283$

GARCHND x = log_art_growth kappa = 50% annual

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
omega	0.0479	0.0291	1.6458	0.0998	*
alpha	0.0800	0.0107	7.5076	0.0000	***
beta	0.9176	0.0160	57.3023	0.0000	***
gamma	-1.3931	0.2035	-6.8442	0.0000	***
nu	11.3272	2.3644	4.7907	0.0000	***

Log-likelihood : -5951.2465
AIC : 11912.4931
BIC : 11942.1171
alpha + beta : 0.9976 (stationary (<1))
d3 active in-sample ($\sigma^2_{t-1} \geq \kappa$): 0.063
Estimated degrees of freedom: $\nu_{\text{hat}} = 11.327$

7. EGARCH(1,1)-X — Extension model

Conditional variance:

$$\ln(\sigma_t^2) = \omega + \alpha [|z_{t-1}| - E|z_{t-1}|] + \xi \cdot z_{t-1} + \beta \cdot \ln(\sigma_{t-1}^2) + \gamma_1 \cdot \text{tone}_{t-1} + \gamma_2 \cdot \text{log_art_growth}_{t-1}$$

What we want to demonstrate: Log-variance specification (Nelson, 1991) with two exogenous regressors. It (i) guarantees positivity of σ^2 without parameter restrictions, (ii) captures asymmetry through the leverage term ξ (negative shocks weigh more than positive ones), and (iii) allows tone and log_art_growth to enter linearly (not as squares). We want to show that EGARCH-X captures asymmetry more cleanly than GARCHAND, and to test whether γ_1 and γ_2 are significant in log-variance rather than in squared variance.

Gaussian

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
mu	-0.1007	0.0375	-2.6841	0.0073	***
omega	0.0342	0.0146	2.3521	0.0187	**
alpha	0.3033	0.0625	4.8542	0.0000	***
xi	-0.0197	0.0168	-1.1735	0.2406	
beta	0.9724	0.0139	69.8722	0.0000	***
gamma_tone	-0.0007	0.0098	-0.0685	0.9454	
gamma_log_artg	0.0843	0.0683	1.2342	0.2171	

Log-likelihood : -6011.7360

AIC : 12037.4719

BIC : 12078.9455

beta = 0.9724 (stationary (|beta|<1))

Half-life of a log-variance shock: 24.8 trading days

Student-t

Parameter	Estimate	Std. Err.	t-ratio	p-value	Sig.
mu	-0.0177	0.0025	-6.9528	0.0000	***
omega	0.0319	0.0102	3.1320	0.0017	***
alpha	0.1565	0.0253	6.1860	0.0000	***
xi	-0.0143	0.0107	-1.3360	0.1816	
beta	0.9792	0.0065	150.4773	0.0000	***
gamma_tone	0.0106	0.0049	2.1500	0.0316	**
gamma_log_artg	-0.0099	0.0601	-0.1645	0.8693	
nu	10.7946	2.1052	5.1275	0.0000	***

Log-likelihood : -5951.4540

AIC : 11918.9080

BIC : 11966.3063

beta = 0.9792 (stationary (|beta|<1))

Half-life of a log-variance shock: 33.0 trading days

$\nu = 10.7946$ (Student-t dof; lower $\rightarrow$ heavier tails. Note: the t-ratio for ν tests $H_0: \nu = 0$, which is not the meaningful null. To test Gaussian vs Student-t use an LR test against the Gaussian EGARCH-X.)

Part II — Robustness Checks (Notebooks 06 & 07)

Six diagnostic checks applied to the GARCHND $\kappa=50\%$ Student-t models (the best in-sample specifications). Checks 1–4 are from notebook 06; Checks 5–6 from notebook 07. All checks point to the same conclusion: the in-sample significance of γ is a regime artefact with no out-of-sample value, and the sentiment variable was measuring global chemistry noise rather than REMX-specific information.

Check 1 — Out-of-sample forecasting (h = 1, 5, 22 days)

In-sample fit through 2025-03-31 (2,514 obs). OOS window: 2025-04-01 to 2026-03-31 (251 obs). Baseline is GARCH(1,1) Student-t. DM > 0 means the candidate beats the baseline; p-values are two-sided.

Loss functions (mean over OOS)

Model	h	QLIKE	MSE
GARCH(1,1) Student-t (base)	1	2.4225	48.0561
GARCH(1,1) Student-t (base)	5	2.4561	38.7524
GARCH(1,1) Student-t (base)	22	2.3900	29.7071
GARCHND tone $\kappa=50\%$	1	2.4397	49.2963
GARCHND tone $\kappa=50\%$	5	2.4792	41.8515
GARCHND tone $\kappa=50\%$	22	2.4417	39.1974
GARCHND log_artg $\kappa=50\%$	1	2.4167	43.2401
GARCHND log_artg $\kappa=50\%$	5	2.4561	35.6411
GARCHND log_artg $\kappa=50\%$	22	2.4229	33.9900

Diebold-Mariano test (baseline vs candidate)

Candidate	h	DM-QLIKE	p	DM-MSE	p
GARCHND tone $\kappa=50\%$	1	−6.262	0.0000	−3.150	0.0016
GARCHND tone $\kappa=50\%$	5	−4.882	0.0000	−4.013	0.0001
GARCHND tone $\kappa=50\%$	22	−6.267	0.0000	−3.409	0.0007
GARCHND log_artg $\kappa=50\%$	1	+1.287	0.1982	+5.153	0.0000
GARCHND log_artg $\kappa=50\%$	5	+0.003	0.9975	+1.831	0.0671
GARCHND log_artg $\kappa=50\%$	22	−2.496	0.0125	−1.969	0.0490

Verdict: neither GARCHND specification improves OOS forecasting. GARCHND tone $\kappa=50\%$ is dominated by the baseline across all horizons and loss functions. GARCHND log_artg $\kappa=50\%$ shows marginal MSE gains at h=1 but no QLIKE gain.

Check 2 — Placebo / regime-only test

Two competing specifications are estimated for each GARCHND $\kappa=50\%$ model: (a) the regime-only model $\gamma \cdot d_3$ (no x^2 term), and (b) five placebos replacing the real x_{t-1} with independent random permutations of it. If the original model does not beat these, the γ estimate reflects regime mean-reversion, not news.

x = tone

Specification	LL	AIC	BIC	γ	SE(γ)	t(γ)
original	-5951.97	11913.94	11943.57	-0.8262	0.0602	-13.72
regime-only ($\gamma \cdot d_3$)	-5950.18	11910.37	11939.99	-1.1218	0.0172	-65.10
placebo seed=2026	-5955.76	11921.52	11951.15	-0.1501	0.0283	-5.31
placebo seed=2027	-5949.09	11908.18	11937.80	-0.7548	0.0233	-32.41
placebo seed=2028	-5952.23	11914.46	11944.08	-0.3636	0.0213	-17.03
placebo seed=2029	-5950.81	11911.63	11941.25	-0.8354	0.0175	-47.80
placebo seed=2030	-5952.38	11914.75	11944.38	-0.3935	0.0610	-6.45

Verdict: original does NOT beat the regime-only / placebo baselines (ΔLL vs regime-only = -1.79 ; placebo LL max = -5949.09 vs original -5951.97). News effect is likely a regime artefact.

x = log_art_growth

Specification	LL	AIC	BIC	γ	SE(γ)	t(γ)
original	-5951.25	11912.49	11942.12	-1.4116	0.2184	-6.46
regime-only ($\gamma \cdot d_3$)	-5950.18	11910.37	11939.99	-1.1218	0.0172	-65.10
placebo seed=2026	-5955.15	11920.29	11949.92	-0.8036	0.1258	-6.39
placebo seed=2027	-5955.34	11920.67	11950.30	-0.2475	0.1213	-2.04
placebo seed=2028	-5955.10	11920.19	11949.82	-0.3646	0.1288	-2.83
placebo seed=2029	-5949.72	11909.44	11939.06	-1.0843	0.0814	-13.32
placebo seed=2030	-5954.52	11919.04	11948.66	-0.4900	0.0936	-5.24

Verdict: original does NOT beat the regime-only / placebo baselines (ΔLL vs regime-only = -1.06 ; placebo LL max = -5949.72 vs original -5951.25). News effect is likely a regime artefact.

Check 3 — SE stability across sigmoid sharpness K

Both GARCHND $\kappa=50\%$ models are refitted at $K \in \{5, 10, 20, 50\}$ and with a hard binary indicator ($K \rightarrow \infty$). Stability is measured by the coefficient of variation of γ across K values; $CV < 0.2$ is considered stable.

x = tone (CV = 0.412 — sensitive to K)

K	LL	γ	SE(γ)	t(γ)	converged
K=5	-5952.63	-0.3626	0.0737	-4.92	True
K=10	-5952.42	-0.3769	0.0740	-5.09	True
K=20	-5951.97	-0.8262	0.0602	-13.72	True
K=50	-5952.16	-0.3880	0.0008	-468.3	True
hard d3	-5953.24	-0.3191	0.0989	-3.23	True

x = log_art_growth (CV = 0.119 — stable)

K	LL	γ	SE(γ)	t(γ)	converged
K=5	-5951.98	-1.3781	0.2739	-5.03	True
K=10	-5952.01	-1.0979	0.2493	-4.40	True
K=20	-5951.25	-1.4116	0.2184	-6.46	True
K=50	-5951.42	-1.1276	0.0405	-27.81	True
hard d3	-5951.52	-1.0767	0.0029	-370.8	True

Verdict: the tone specification is sensitive to the sigmoid shape; log_art_growth is stable in sign and magnitude but not in t-ratio (t ranges from -4.4 to -370.8).

Check 4 — High-volatility episode clusters ($\kappa=50\%$)

Days where $\sigma_{t-1}^2 \geq \kappa_{\text{high}}$ (annualised 50% vol) are grouped into episodes separated by gaps of at most 5 days. The longest episode determines the recommended block-bootstrap length and HAC bandwidth.

x = tone: 193 triggered days (6.98%), 21 distinct episodes; recommended block length $\geq$ 50 days

#	Start	End	Days
1	2021-01-25	2021-04-06	50
2	2026-02-03	2026-03-31	40
3	2020-03-04	2020-04-22	35
4	2025-10-14	2025-11-14	24
5	2025-04-11	2025-04-28	11
6	2018-12-24	2019-01-08	10
7	2024-09-30	2024-10-11	10
8	2022-05-11	2022-05-23	9
9	2022-11-08	2022-11-18	9
10	2022-06-03	2022-06-13	7

x = log_art_growth: 179 triggered days (6.47%), 20 distinct episodes; recommended block length $\geq$ 42 days

#	Start	End	Days
1	2021-02-04	2021-04-06	42
2	2020-03-04	2020-04-29	40
3	2025-10-15	2025-11-12	21
4	2026-03-05	2026-03-31	19
5	2021-01-07	2021-01-25	12
6	2026-02-03	2026-02-19	12
7	2024-09-30	2024-10-14	11
8	2018-12-24	2019-01-08	10
9	2022-05-11	2022-05-23	9
10	2022-06-03	2022-06-15	9

Check 5 — HAC and Moving-Block-Bootstrap SEs on γ (Notebook 07)

Sandwich (HAC) and moving-block-bootstrap (MBB, B=2000) standard errors on γ for both GARCHND $\kappa=50\%$ specifications, at block lengths $L \in \{25, 50, 75, 100\}$ days (informed by the episode clusters in Check 4).

x = tone

SE method	γ	SE(γ)	t(γ)
iid sandwich	-0.8262	0.0602	-13.72
HAC L=25	-0.8262	0.0586	-14.09
HAC L=50	-0.8262	0.0548	-15.07
HAC L=75	-0.8262	0.0565	-14.62
HAC L=100	-0.8262	0.0578	-14.30
MBB L=25	-0.8262	0.0583	-14.18
MBB L=50	-0.8262	0.0563	-14.68
MBB L=75	-0.8262	0.0575	-14.36
MBB L=100	-0.8262	0.0596	-13.87

x = log_art_growth

SE method	γ	SE(γ)	t(γ)
iid sandwich	-1.4116	0.2184	-6.46
HAC L=25	-1.4116	0.2321	-6.08
HAC L=50	-1.4116	0.2424	-5.82
HAC L=75	-1.4116	0.2469	-5.72
HAC L=100	-1.4116	0.2473	-5.71
MBB L=25	-1.4116	0.2360	-5.98
MBB L=50	-1.4116	0.2455	-5.75
MBB L=75	-1.4116	0.2471	-5.71
MBB L=100	-1.4116	0.2503	-5.64

Verdict: HAC and MBB corrections do NOT inflate the standard errors materially. The t-ratios remain strongly significant (tone: around -14 ; log_artg: around -5.7). This is explained by the iid nature of x_{t-1}^2 within the regime, which absorbs serial correlation in the score. The conclusion from Checks 1–2 (regime artefact, no OOS gain) stands unaffected.

Check 6 — GDELT coverage of REMX / rare earths (Notebook 07)

The GDELT article cache (314,048 articles, 2015-04-01 to 2026-04-01) was built with a permissive 19-concept filter (any single match passes). This check quantifies how many articles are genuinely related to the REMX economy across four progressively tighter filter tiers.

Articles per concept (full cache)

Concept	N articles	Share
has_cerium	186,346	59.34%
has_rare_earth	92,383	29.42%
has_lynas	26,066	8.30%
has_samarium	5,005	1.59%
has_neodymium	2,960	0.94%
has_mp_materials	2,891	0.92%
has_praseodymium	2,124	0.68%
has_arafura	1,730	0.55%
has_europium	1,153	0.37%
has_terbium	1,143	0.36%
has_gadolinium	1,051	0.33%
has_yttrium	1,030	0.33%
has_lanthanum	877	0.28%
has_china_northern	702	0.22%
has_dysprosium	576	0.18%
has_ndfeb	505	0.16%
has_ndpr	454	0.14%
has_remx	150	0.05%
has_jl_mag	28	0.01%

Filter tiers

Tier	N articles	Art/day	% zero-days	Mean tone	Std tone
Broad (≥ 1 concept, used in model)	314,048	78.5	0.5%	-0.280	3.299
Medium (metal OR company)	231,708	57.9	0.5%	-0.453	3.585
Tight (rare_earth AND metal/co.)	10,158	2.5	38.5%	-0.130	2.133
Ultra-tight (metal AND company)	559	0.14	91.4%	+0.638	1.850